

Matrix Multiplication Optimization: Comparison of Dense and Sparse Algorithms

Contents

| | | |
|----------|--|----------|
| 1 | Introduction | 2 |
| 2 | Methodology | 2 |
| 2.1 | Testing Platform | 2 |
| 2.2 | Test Configuration | 2 |
| 3 | Results: Dense Matrices | 2 |
| 3.1 | Execution Time Analysis | 2 |
| 3.2 | Memory Usage | 3 |
| 3.3 | CPU Usage | 3 |
| 3.4 | Performance Analysis | 3 |
| 4 | Results: Sparse Matrices | 4 |
| 4.1 | Effect of Sparsity Level | 4 |
| 4.2 | Sparse vs. Dense Comparison | 5 |
| 5 | Performance Bottlenecks | 6 |
| 5.1 | Dense Matrices | 6 |
| 5.2 | Sparse Matrices | 6 |
| 6 | Maximum Matrix Size Handled | 7 |
| 7 | Conclusions and Recommendations | 7 |
| 7.1 | Key Conclusions | 7 |
| 7.2 | Practical Recommendations | 7 |

1 Introduction

Matrix multiplication is a fundamental operation in scientific computing, machine learning, and computer graphics. The basic algorithm has $O(n^3)$ complexity, making it computationally expensive for large matrices. This work investigates different optimization strategies:

- **Basic algorithm (ijk):** Standard implementation with three nested loops
- **Cache Optimized (ikj):** Loop reordering to improve cache locality
- **Blocked multiplication:** Division of matrices into smaller blocks
- **Sparse matrices:** Specialized data structures for matrices with many zeros

2 Methodology

2.1 Testing Platform

Experiments were executed in Python using the following metrics:

- **Execution time:** Measured with `time.time()`
- **Memory:** Measured with `tracemalloc` (peak memory usage)
- **CPU:** Monitored with `psutil`

2.2 Test Configuration

- **Matrix sizes (dense):** 128×128 , 256×256 , 512×512
- **Matrix size (sparse):** 500×500
- **Sparsity levels:** 70%, 90%, 95%, 99%
- **Block size:** 32 for blocked multiplication

3 Results: Dense Matrices

3.1 Execution Time Analysis

Table 1 shows the execution times for different matrix sizes and algorithms.

Table 1: Execution time for dense matrices (seconds)

| Size | Basic (ijk) | Cache Opt (ikj) | Blocked (32) |
|------------------|-------------|-----------------|--------------|
| 128×128 | 2.6089 | 4.1018 | 6.8700 |
| 256×256 | 28.8605 | 31.8817 | 54.7863 |
| 512×512 | 368.4366 | 373.7061 | 645.5046 |




Figure 1: Performance comparison of dense matrix multiplication algorithms. Left: Execution time vs matrix size showing exponential growth. Center: Memory usage increases quadratically with matrix size. Right: Speedup comparison relative to the basic algorithm baseline.

Figure 1 illustrates the performance characteristics of the three dense matrix multiplication algorithms across different matrix sizes. The execution time grows rapidly with matrix size, following the expected $O(n^3)$ complexity.

3.2 Memory Usage

Table 2 presents memory consumption for each method.

Table 2: Memory usage for dense matrices (MB)

| Size | Basic (ijk) | Cache Opt (ijk) | Blocked (32) |
|---------|-------------|-----------------|--------------|
| 128×128 | 0.51 | 0.51 | 0.51 |
| 256×256 | 4.74 | 2.04 | 2.04 |
| 512×512 | 8.07 | 8.07 | 8.07 |

3.3 CPU Usage

CPU utilization remained consistently high across all methods (94-97%), indicating compute-bound operations.

3.4 Performance Analysis

Key Observations:

- For the tested configuration, the **basic algorithm** was fastest
- Cache optimized and blocked methods showed **slower performance** than expected
- This suggests that for smaller matrices on this specific hardware, the basic approach has less overhead
- Memory usage is similar across all methods (~0.5-8 MB)

Speedup Analysis (relative to Basic):

- Cache Optimized: 0.64x - 0.99x (slower in most cases)

- Blocked (32): 0.38x - 0.57x (consistently slower)

Possible Explanations:

1. Python's overhead may dominate for these matrix sizes
2. Cache optimizations may benefit from larger matrices (>1024)
3. Block size (32) may not be optimal for this hardware
4. Compiled languages (C/C++) would show more significant differences

4 Results: Sparse Matrices

4.1 Effect of Sparsity Level

Table 3 shows how sparsity level affects performance for 500×500 matrices.

Table 3: Performance vs. Sparsity Level (500×500 matrices)

| Sparsity | Time (s) | Memory (MB) | CPU % | Speedup |
|----------|----------|-------------|-------|-----------------|
| 70% | 812.4857 | 29.08 | 95.8 | 1.00x |
| 90% | 74.3941 | 28.95 | 96.4 | 10.92x |
| 95% | 12.2341 | 28.34 | 97.2 | 66.41x |
| 99% | 0.4434 | 1.68 | 103.5 | 1832.52x |




Figure 2: Impact of sparsity on matrix multiplication performance. Top left: Execution time decreases exponentially with sparsity level (logarithmic scale). Top right: Memory usage remains stable until very high sparsity (99%). Bottom left: Performance improvement relative to 70% sparsity baseline, showing dramatic 1824.6x speedup at 99% sparsity. Bottom right: Direct comparison between sparse and dense methods at 90% sparsity, demonstrating 7.83x faster execution and reduced memory footprint.

Key Findings:

- **Dramatic improvement** with high sparsity: 99% sparsity is 1832x faster than 70%
- Memory reduction: 94.2% less memory at 99% vs 70% sparsity
- Sparse matrices are highly efficient when >90% of elements are zero

As shown in Figure 2, the relationship between sparsity and performance is not linear. The most significant gains occur at very high sparsity levels (95% and above), where the sparse representation can skip the vast majority of computations.

4.2 Sparse vs. Dense Comparison

Table 4 compares sparse and dense methods at 90% sparsity.

Table 4: Sparse vs. Dense multiplication (500×500 , 90% sparsity)

| Method | Time (s) | Memory (MB) | Speedup |
|--------|----------|-------------|--------------|
| Dense | 450.9265 | 38.62 | 1.00x |
| Sparse | 53.2125 | 28.95 | 8.47x |

Results:

- Sparse method is **8.47x faster** than dense at 90% sparsity
- Memory savings: 25.0% reduction
- Clear advantage for sparse representation when most elements are zero

5 Performance Bottlenecks

5.1 Dense Matrices

Basic Algorithm (ijk):

- Memory access pattern: Non-sequential access to matrix B
- Cache misses: High due to column-wise access
- Scalability: $O(n^3)$ - grows cubically

Cache Optimized (ikj):

- Improvement: Better spatial locality
- Limitation: Python overhead may mask benefits for small matrices
- Expected benefit: More significant for matrices $> 1024 \times 1024$

Blocked Multiplication:

- Improvement: Blocks fit in L1/L2 cache
- Limitation: Control overhead for small matrices
- Optimal block size: Architecture-dependent (typically 32-128)

5.2 Sparse Matrices

Advantages:

- Only computes non-zero elements
- Reduced memory Fetch
- Ideal for sparsity $> 90\%$

Limitations:

- Data structure overhead for low sparsity ($< 70\%$)
- Dictionary/map lookup costs in Python
- Less efficient when many elements are non-zero

6 Maximum Matrix Size Handled

Based on execution times and memory constraints:

Table 5: Maximum practical matrix sizes

| Method | Max Size (approximate) |
|-----------------------|------------------------|
| Basic (ijk) | 512×512 (6 min) |
| Cache Optimized | 512×512 (6 min) |
| Blocked (32) | 512×512 (11 min) |
| Sparse (90% sparsity) | 1000×1000+ |
| Sparse (99% sparsity) | 2000×2000+ |

Note: Maximum sizes are limited by reasonable execution time (<15 minutes) rather than memory constraints.

7 Conclusions and Recommendations

7.1 Key Conclusions

1. **Dense matrices (small to medium):** The basic algorithm performed adequately for matrices up to 512×512 in this Python implementation.
2. **Cache optimizations:** Did not show expected improvements in this configuration, likely due to:
 - Python's interpreted nature and overhead
 - Matrix sizes still relatively small
 - Would benefit more in compiled languages (C/C++, Java)
3. **Sparse matrices:** Provide **dramatic improvements** when sparsity > 90%:
 - Up to 1832x faster at 99% sparsity
 - 8.47x faster than dense at 90% sparsity
 - Significant memory savings
4. **Sparsity is the key factor:** The level of sparsity has more impact on performance than cache optimizations for this implementation.

7.2 Practical Recommendations

For Dense Matrices:

- Small matrices (<256): Use basic algorithm in Python
- Large matrices (>512): Consider NumPy, BLAS libraries, or compiled languages
- Cache optimizations: More effective in C/C++/Java implementations

For Sparse Matrices:

- **Always use sparse** when sparsity > 90%
- **Consider sparse** when sparsity is 70-90% (test both)
- **Use dense** when sparsity < 70%

For Production Systems:

- Use specialized libraries: NumPy, SciPy (sparse), BLAS, cuBLAS (GPU)
- Profile before optimizing: Measure actual bottlenecks
- Consider hardware: GPU acceleration for large matrices